

针对问题:

在十一号机子上，数据读取的时间占了总算法运行时间的一半甚至更多，在提升 dataloader 的 num_workers 同时并没有改善。数据读取占用很久时间的现象理应是不正常的，同时在监控 GPU 运行效率是会发现 GPU 会浪费很多时间在等待读取完的数据输入，即 GPU 利用率会间隔一定时间持续为；同时，这种现象在小规模数据集上不会对总体运行时间带来很大的影响，但在大规模数据集上的每个 batch 运行时间的累加会带来非常大的影响。

怀疑因素：

1. 11 号服务器上的硬盘不是直连的，通过 nas 连接，在每次读取 batch 数据的时候，I/O 口的性能瓶颈；
2. 数据处理的代码段消耗太多时间。

实验：

在 11, 7, 8 号机子上做了相同的读取数据实验，并分别记录了硬盘读取和数据预处理的时间，并在不同的 num_workers(0, 1, 2, 4) 的条件下进行实验。实际测算结果如下表所示：

服务器	Batchsize: 64											
	0			1			2			4		
	总数据 时间	硬盘读 取时间	数据处 理时间	总数据 时间	硬盘读 取时间	数据处 理时间	总数据 时间	硬盘读 取时间	数据处 理时间	总数据 时间	硬盘读 取时间	数据处 理时间
11 号机	4.514	0.223	3.048	0.2699	0.0737	0.1419	0.2042	0.0823	0.1535	0.1591	0.0876	0.1662
7 号机	0.2602	0.0532	0.1732	0.3010	0.0550	0.1800	0.2438	0.0617	0.2176	0.2525	0.0654	0.2016
8 号机	0.1196	0.0268	0.0778	0.1493	0.0280	0.0917	0.1345	0.0654	0.1026	0.1212	0.0350	0.0916

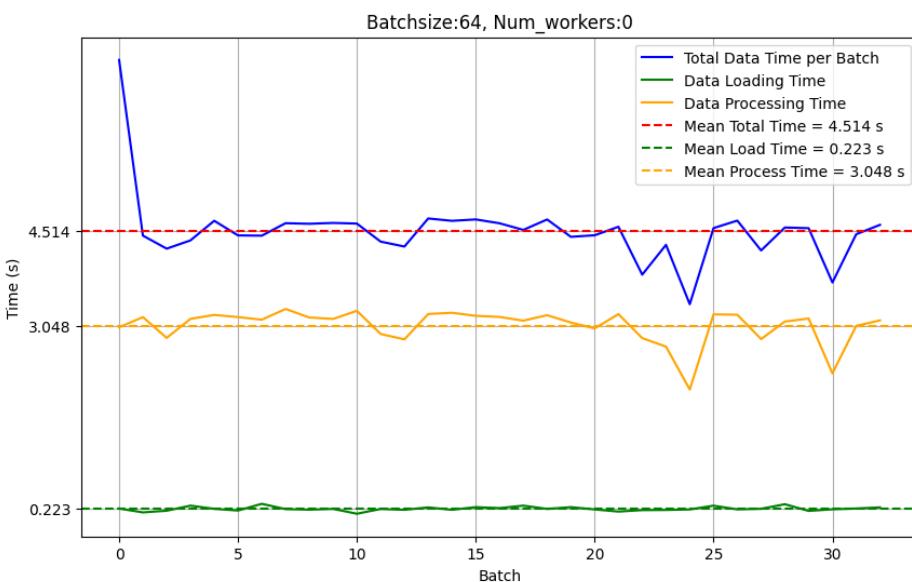
可以看到，8 号机子上的总体偏优，基本比其它两个机子提升两倍之多，我不太了解各服务器的配置，所以暂时不知道具体原因。

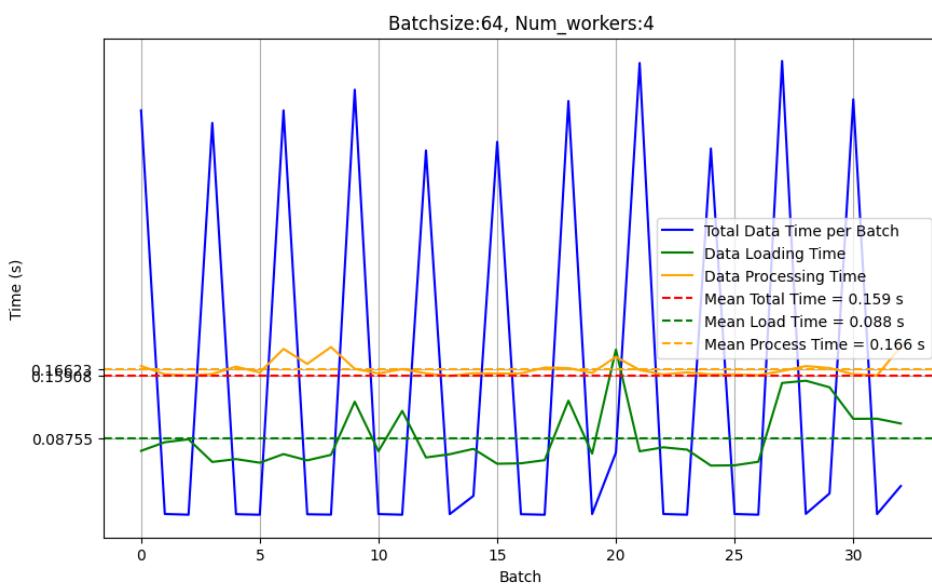
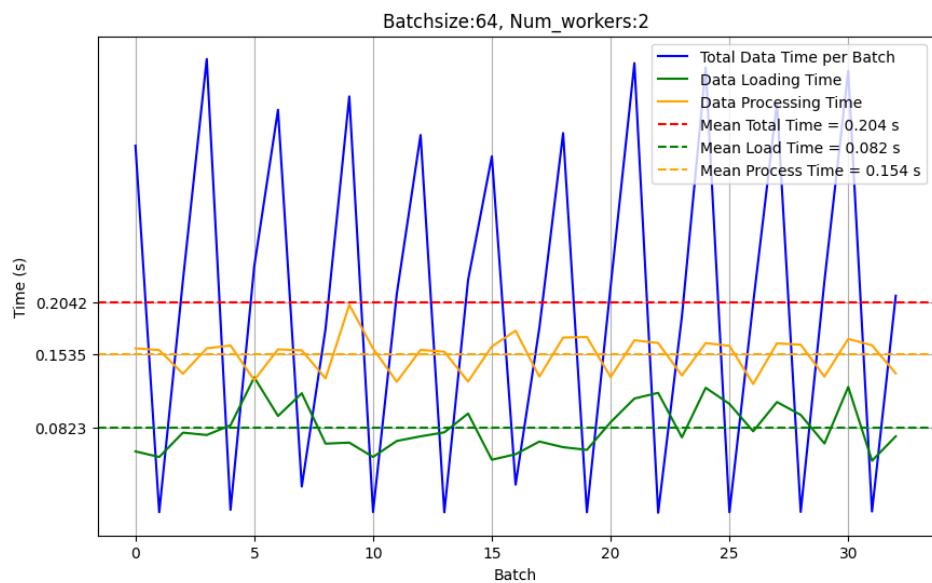
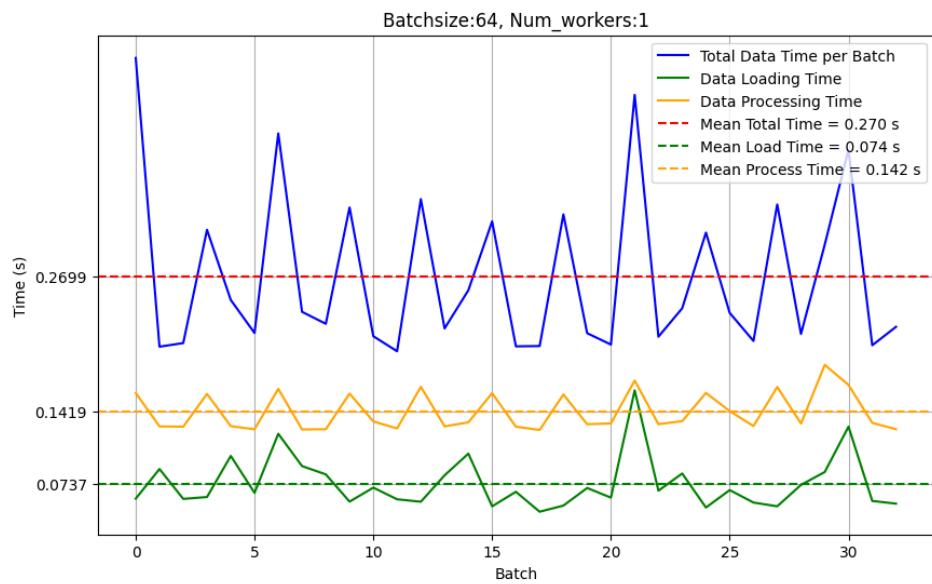
暂时能分析出来：

- 1 同样的数据处理段为什么差别这么大，是因为不同 cpu 的差异吗，毕竟差两倍，在大规模数据集上是 10 小时和 5 小时的区别。

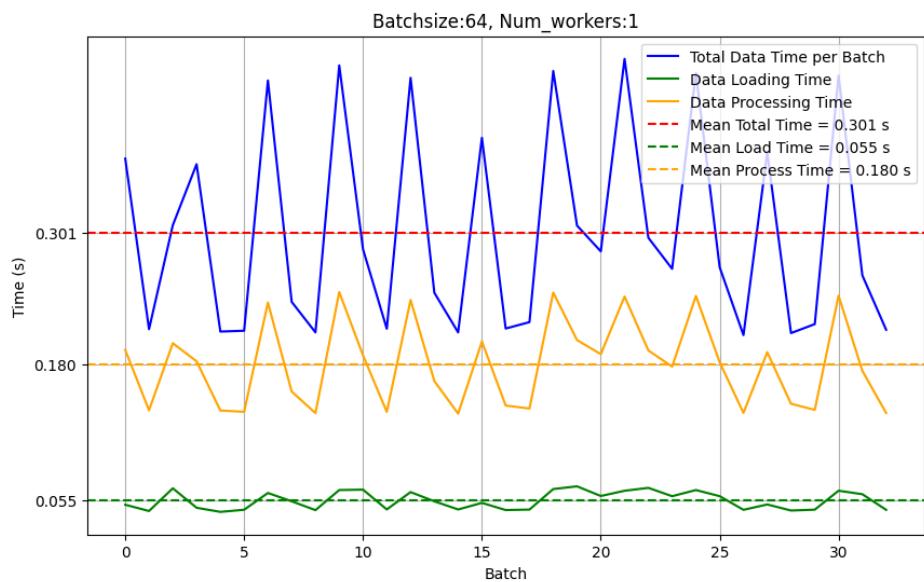
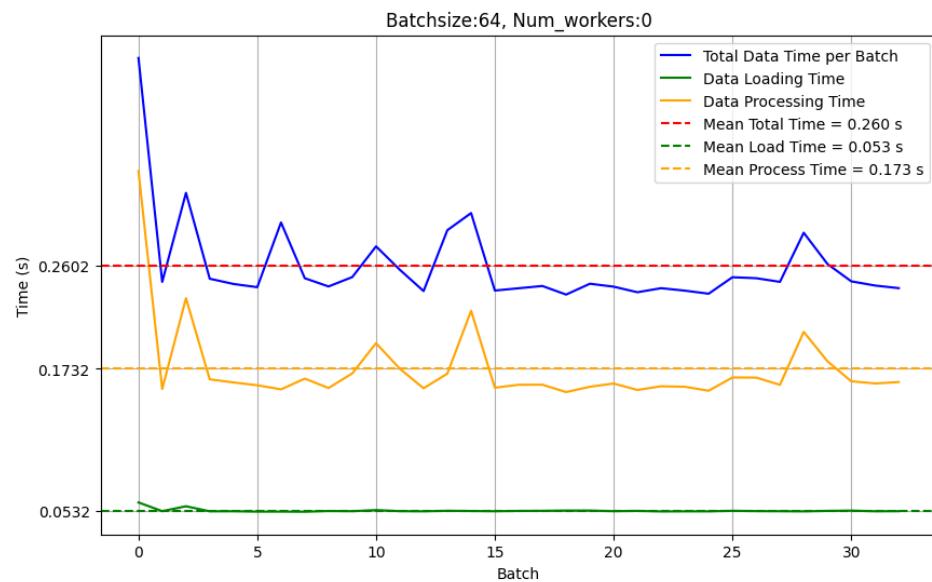
如下是实验截图：

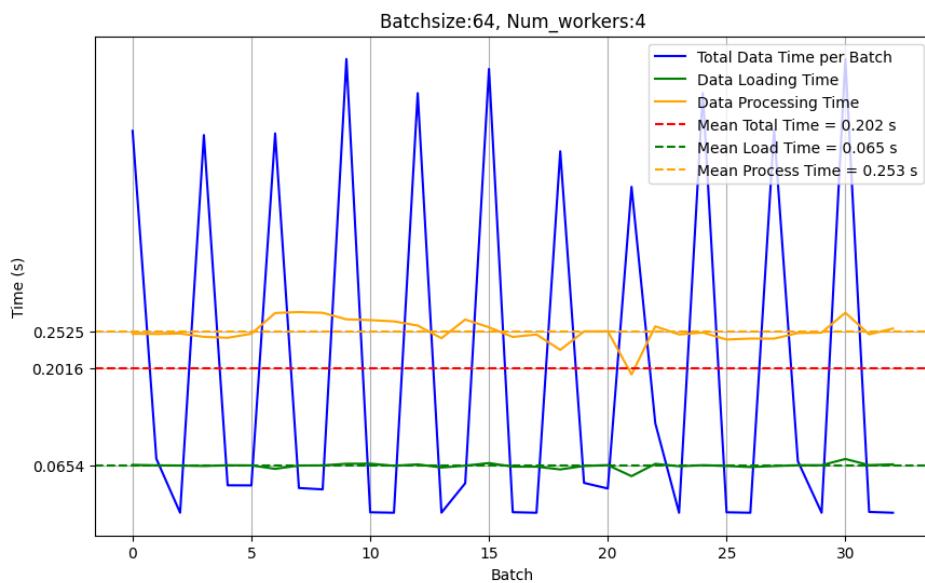
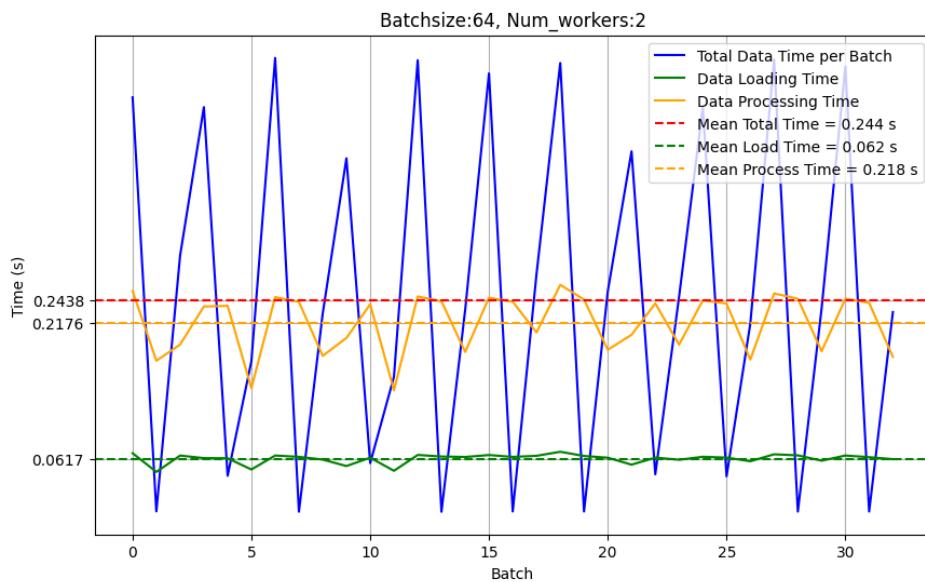
11 号机：



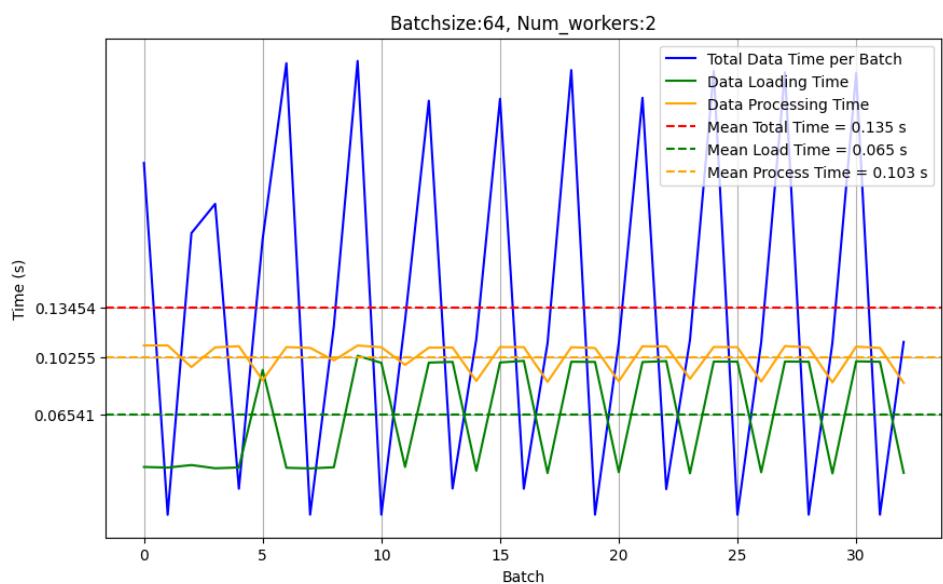
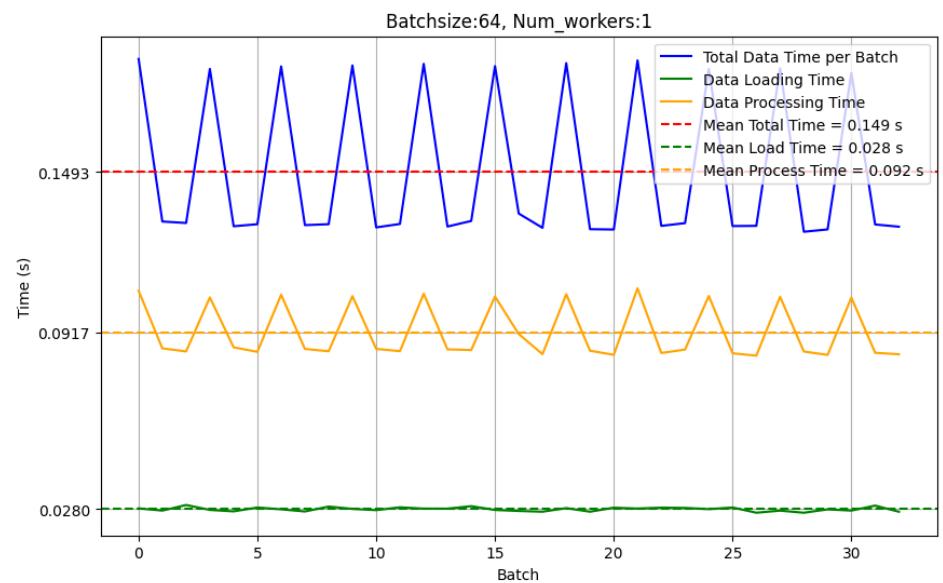
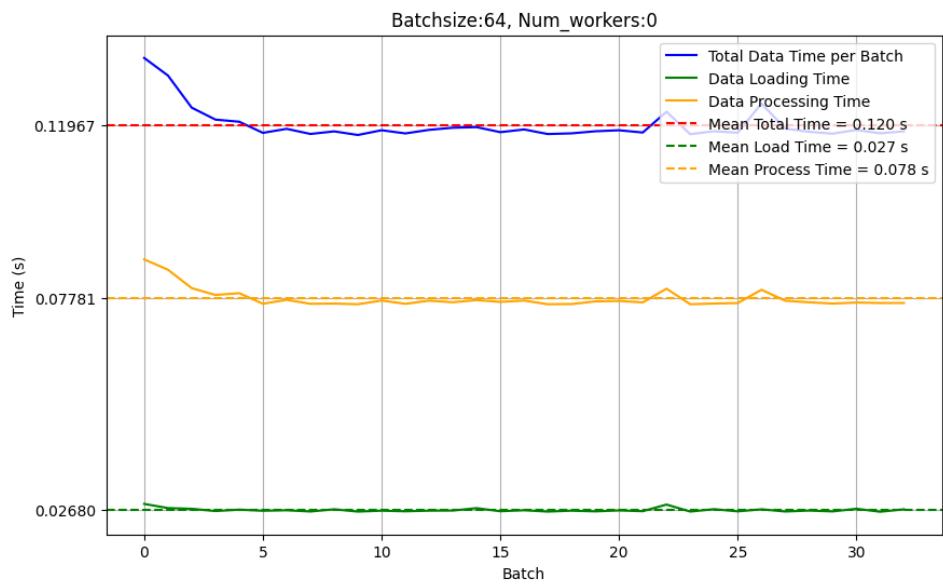


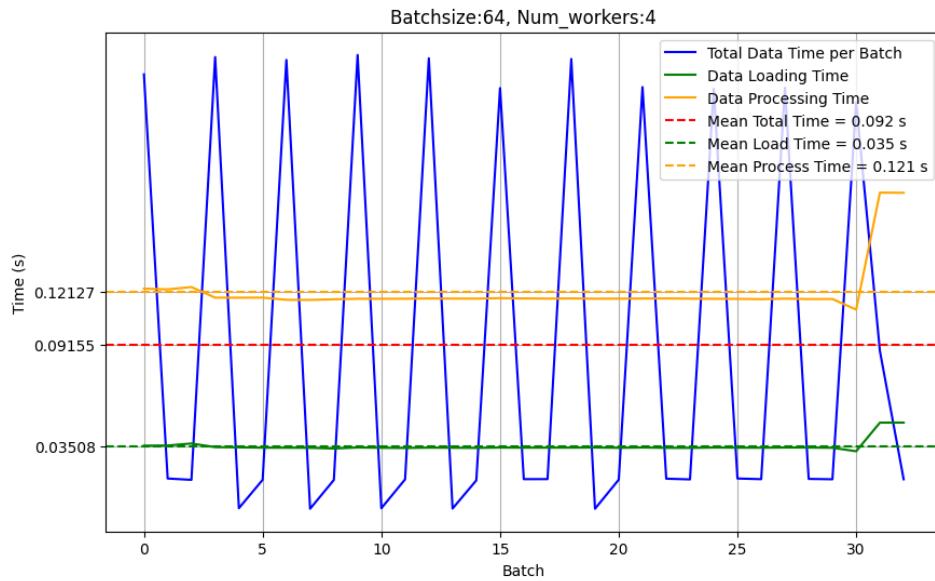
7号机:





8号机:





实际运行：

Num_workers=4:

```
- Using Data parallel ...
- [Epoch 0/200][Batch 1/11498] BatchTime = 17.118 (s) DataTime = 12.201 (s) Losses =
- [Epoch 0/200][Batch 2/11498] BatchTime = 0.897 (s) DataTime = 0.017 (s) Losses =
- [Epoch 0/200][Batch 3/11498] BatchTime = 0.888 (s) DataTime = 0.006 (s) Losses =
- [Epoch 0/200][Batch 4/11498] BatchTime = 0.888 (s) DataTime = 0.007 (s) Losses =
- [Epoch 0/200][Batch 5/11498] BatchTime = 4.048 (s) DataTime = 3.150 (s) Losses =
- [Epoch 0/200][Batch 6/11498] BatchTime = 1.950 (s) DataTime = 1.065 (s) Losses =
- [Epoch 0/200][Batch 7/11498] BatchTime = 0.887 (s) DataTime = 0.007 (s) Losses =
- [Epoch 0/200][Batch 8/11498] BatchTime = 0.893 (s) DataTime = 0.009 (s) Losses =
- [Epoch 0/200][Batch 9/11498] BatchTime = 7.183 (s) DataTime = 6.283 (s) Losses =
- [Epoch 0/200][Batch 10/11498] BatchTime = 1.202 (s) DataTime = 0.325 (s) Losses =
- [Epoch 0/200][Batch 11/11498] BatchTime = 0.889 (s) DataTime = 0.007 (s) Losses =
- [Epoch 0/200][Batch 12/11498] BatchTime = 0.892 (s) DataTime = 0.017 (s) Losses =
- [Epoch 0/200][Batch 13/11498] BatchTime = 6.075 (s) DataTime = 5.179 (s) Losses =
- [Epoch 0/200][Batch 14/11498] BatchTime = 1.232 (s) DataTime = 0.350 (s) Losses =
- [Epoch 0/200][Batch 15/11498] BatchTime = 0.892 (s) DataTime = 0.008 (s) Losses =
- [Epoch 0/200][Batch 16/11498] BatchTime = 0.897 (s) DataTime = 0.008 (s) Losses =
- [Epoch 0/200][Batch 17/11498] BatchTime = 8.337 (s) DataTime = 7.447 (s) Losses =
- [Epoch 0/200][Batch 18/11498] BatchTime = 1.049 (s) DataTime = 0.167 (s) Losses =
- [Epoch 0/200][Batch 19/11498] BatchTime = 0.891 (s) DataTime = 0.008 (s) Losses =
- [Epoch 0/200][Batch 20/11498] BatchTime = 0.889 (s) DataTime = 0.010 (s) Losses =
- [Epoch 0/200][Batch 21/11498] BatchTime = 7.363 (s) DataTime = 6.464 (s) Losses =
- [Epoch 0/200][Batch 22/11498] BatchTime = 2.696 (s) DataTime = 1.804 (s) Losses =
- [Epoch 0/200][Batch 23/11498] BatchTime = 0.893 (s) DataTime = 0.008 (s) Losses =
- [Epoch 0/200][Batch 24/11498] BatchTime = 0.902 (s) DataTime = 0.015 (s) Losses =
- [Epoch 0/200][Batch 25/11498] BatchTime = 5.857 (s) DataTime = 4.953 (s) Losses =
- [Epoch 0/200][Batch 26/11498] BatchTime = 2.195 (s) DataTime = 1.308 (s) Losses =
```

Num_works:16

```
[Epoch 0/200][Batch 1/11498] BatchTime = 33.885 (s) DateTime = 28.682 (s) Losses = ['0.0000', '137.9634'] lr = 0.000000
[Epoch 0/200][Batch 2/11498] BatchTime = 0.903 (s) DateTime = 0.015 (s) Losses = ['0.0000', '137.8153'] lr = 0.001000
[Epoch 0/200][Batch 3/11498] BatchTime = 0.894 (s) DateTime = 0.009 (s) Losses = ['0.0000', '197.5318'] lr = 0.001000
[Epoch 0/200][Batch 4/11498] BatchTime = 0.899 (s) DateTime = 0.009 (s) Losses = ['0.0000', '149.4078'] lr = 0.001000
[Epoch 0/200][Batch 5/11498] BatchTime = 6.743 (s) DateTime = 5.849 (s) Losses = ['0.0000', '108.9711'] lr = 0.001000
[Epoch 0/200][Batch 6/11498] BatchTime = 1.724 (s) DateTime = 0.844 (s) Losses = ['0.0000', '104.7835'] lr = 0.001000
[Epoch 0/200][Batch 7/11498] BatchTime = 0.903 (s) DateTime = 0.016 (s) Losses = ['0.0000', '88.8002'] lr = 0.001000
[Epoch 0/200][Batch 8/11498] BatchTime = 0.894 (s) DateTime = 0.008 (s) Losses = ['0.0000', '77.5628'] lr = 0.001000
[Epoch 0/200][Batch 9/11498] BatchTime = 0.890 (s) DateTime = 0.008 (s) Losses = ['0.0000', '79.8887'] lr = 0.001000
[Epoch 0/200][Batch 10/11498] BatchTime = 0.891 (s) DateTime = 0.008 (s) Losses = ['0.0000', '80.2600'] lr = 0.001000
[Epoch 0/200][Batch 11/11498] BatchTime = 0.886 (s) DateTime = 0.007 (s) Losses = ['0.0000', '71.8680'] lr = 0.001000
[Epoch 0/200][Batch 12/11498] BatchTime = 0.887 (s) DateTime = 0.007 (s) Losses = ['0.0000', '71.3163'] lr = 0.001000
[Epoch 0/200][Batch 13/11498] BatchTime = 0.886 (s) DateTime = 0.008 (s) Losses = ['0.0000', '70.3161'] lr = 0.001000
[Epoch 0/200][Batch 14/11498] BatchTime = 0.880 (s) DateTime = 0.007 (s) Losses = ['0.0000', '71.2721'] lr = 0.001000
[Epoch 0/200][Batch 15/11498] BatchTime = 0.885 (s) DateTime = 0.007 (s) Losses = ['0.0000', '70.8956'] lr = 0.001000
[Epoch 0/200][Batch 16/11498] BatchTime = 0.904 (s) DateTime = 0.019 (s) Losses = ['0.0000', '70.8994'] lr = 0.001000
[Epoch 0/200][Batch 17/11498] BatchTime = 10.369 (s) DateTime = 9.479 (s) Losses = ['0.0000', '63.0775'] lr = 0.001000
[Epoch 0/200][Batch 18/11498] BatchTime = 5.434 (s) DateTime = 4.542 (s) Losses = ['0.0000', '63.0981'] lr = 0.001000
[Epoch 0/200][Batch 19/11498] BatchTime = 0.890 (s) DateTime = 0.008 (s) Losses = ['0.0000', '62.1150'] lr = 0.001000
[Epoch 0/200][Batch 20/11498] BatchTime = 0.891 (s) DateTime = 0.008 (s) Losses = ['0.0000', '61.7713'] lr = 0.001000
[Epoch 0/200][Batch 21/11498] BatchTime = 10.413 (s) DateTime = 9.524 (s) Losses = ['0.0000', '65.3981'] lr = 0.001000
[Epoch 0/200][Batch 22/11498] BatchTime = 3.658 (s) DateTime = 2.752 (s) Losses = ['0.0000', '66.5505'] lr = 0.001000
[Epoch 0/200][Batch 23/11498] BatchTime = 0.895 (s) DateTime = 0.010 (s) Losses = ['0.0000', '67.6510'] lr = 0.001000
[Epoch 0/200][Batch 24/11498] BatchTime = 0.887 (s) DateTime = 0.009 (s) Losses = ['0.0000', '65.0815'] lr = 0.001000
[Epoch 0/200][Batch 25/11498] BatchTime = 0.889 (s) DateTime = 0.007 (s) Losses = ['0.0000', '63.0594'] lr = 0.001000
[Epoch 0/200][Batch 26/11498] BatchTime = 0.884 (s) DateTime = 0.008 (s) Losses = ['0.0000', '64.2656'] lr = 0.001000
[Epoch 0/200][Batch 27/11498] BatchTime = 2.786 (s) DateTime = 1.892 (s) Losses = ['0.0000', '63.4628'] lr = 0.001000
[Epoch 0/200][Batch 28/11498] BatchTime = 0.895 (s) DateTime = 0.013 (s) Losses = ['0.0000', '62.3571'] lr = 0.001000
[Epoch 0/200][Batch 29/11498] BatchTime = 0.892 (s) DateTime = 0.008 (s) Losses = ['0.0000', '63.9437'] lr = 0.001000
[Epoch 0/200][Batch 30/11498] BatchTime = 0.887 (s) DateTime = 0.008 (s) Losses = ['0.0000', '61.0530'] lr = 0.001000
[Epoch 0/200][Batch 31/11498] BatchTime = 0.894 (s) DateTime = 0.008 (s) Losses = ['0.0000', '60.0711'] lr = 0.001000
[Epoch 0/200][Batch 32/11498] BatchTime = 0.901 (s) DateTime = 0.014 (s) Losses = ['0.0000', '58.8759'] lr = 0.001000
[Epoch 0/200][Batch 33/11498] BatchTime = 11.617 (s) DateTime = 10.709 (s) Losses = ['0.0000', '63.6782'] lr = 0.001000
[Epoch 0/200][Batch 34/11498] BatchTime = 4.482 (s) DateTime = 3.581 (s) Losses = ['0.0000', '58.2843'] lr = 0.001000
[Epoch 0/200][Batch 35/11498] BatchTime = 0.894 (s) DateTime = 0.009 (s) Losses = ['0.0000', '60.3085'] lr = 0.001000
[Epoch 0/200][Batch 36/11498] BatchTime = 0.889 (s) DateTime = 0.009 (s) Losses = ['0.0000', '64.3093'] lr = 0.001000
[Epoch 0/200][Batch 37/11498] BatchTime = 10.303 (s) DateTime = 9.376 (s) Losses = ['0.0000', '65.6142'] lr = 0.001000
[Epoch 0/200][Batch 38/11498] BatchTime = 2.136 (s) DateTime = 1.236 (s) Losses = ['0.0000', '61.2851'] lr = 0.001000
[Epoch 0/200][Batch 39/11498] BatchTime = 0.918 (s) DateTime = 0.011 (s) Losses = ['0.0000', '53.7125'] lr = 0.001000
[Epoch 0/200][Batch 40/11498] BatchTime = 1.578 (s) DateTime = 0.677 (s) Losses = ['0.0000', '59.6389'] lr = 0.001000
[Epoch 0/200][Batch 41/11498] BatchTime = 0.903 (s) DateTime = 0.011 (s) Losses = ['0.0000', '60.8825'] lr = 0.001000
[Epoch 0/200][Batch 42/11498] BatchTime = 0.884 (s) DateTime = 0.009 (s) Losses = ['0.0000', '66.3812'] lr = 0.001000
[Epoch 0/200][Batch 43/11498] BatchTime = 0.893 (s) DateTime = 0.008 (s) Losses = ['0.0000', '58.0004'] lr = 0.001000
```